

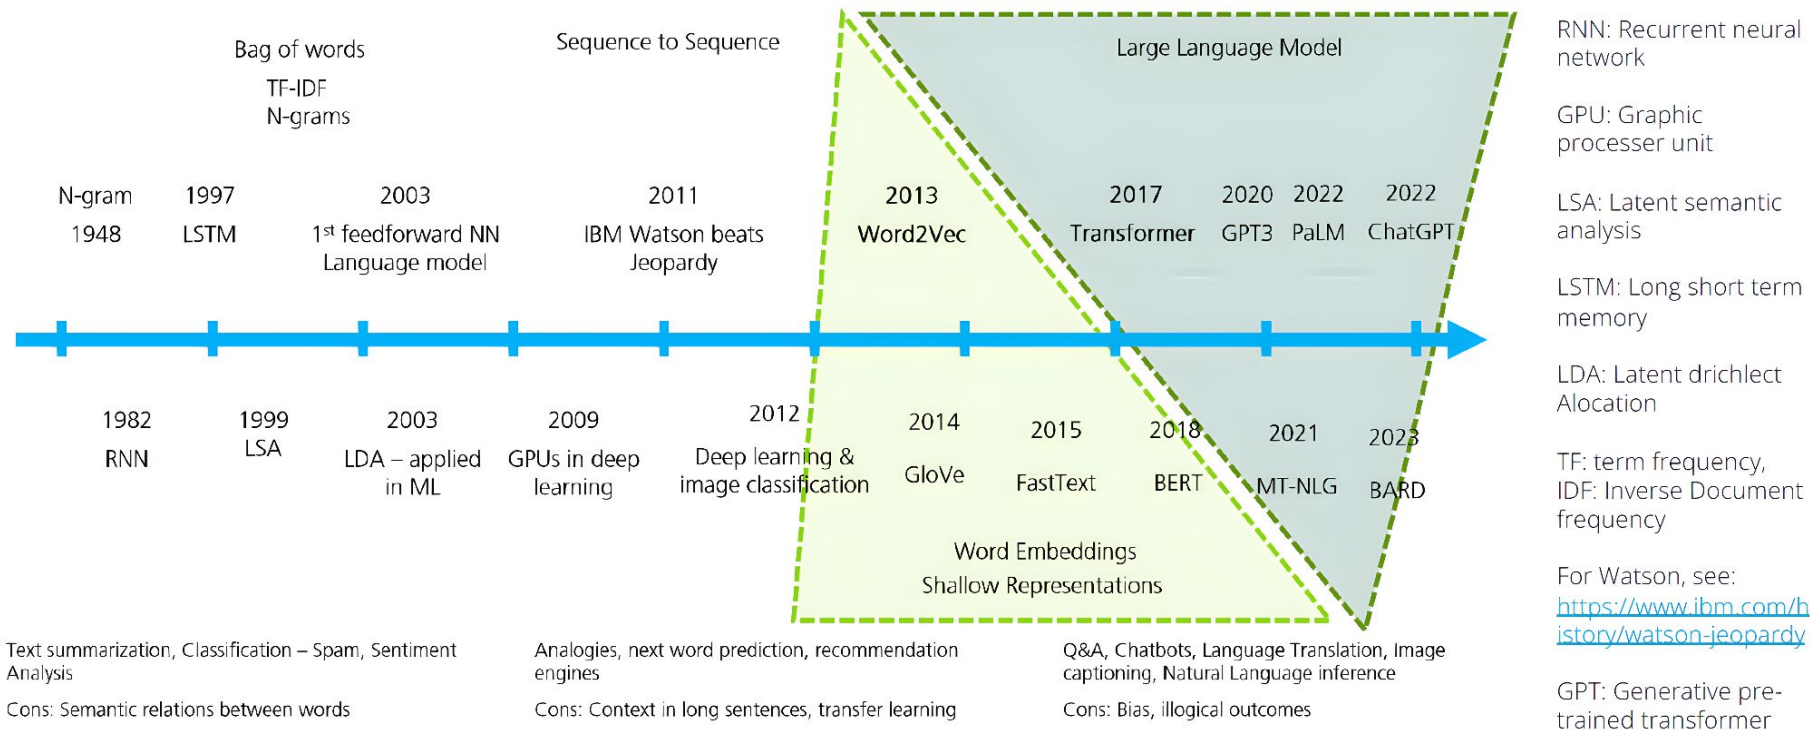
Text Analysis II

Introduction to Computational Social Science

Outline



- From word embeddings to Transformers
 - Encoding and Decoding: Understanding Modern Models
 - Fine-tuning and Adaptation
 - Generative AI and Prompt Engineering
-
- Tutorial: Sentiment and Topic Modeling in Parliamentary Speech
 - Discussion: The Promise and Pitfalls of NLP for Social Sciences

From word embeddings to Transformers




Recurrent Neural Networks and LSTMs

one to one



Recurrent Neural Networks and LSTMs

one to one



one to many






Image Captioning

many to one




Sentiment

many to many




Machine Translation

many to many





Named Entity
Recognition

Recurrent Neural Networks and LSTMs




Issues: slow to train; memory loss

Word Embeddings





$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_t | w_{t+j})$$




$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

Word Embeddings



$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_t | w_{t+j})$$



$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

Issues: aggregates contexts; single embedding for a unique word (e.g., Slack)

Transformers (Attention Mechanism)

Attention is all you need

[PDF] [neurips.cc](#)

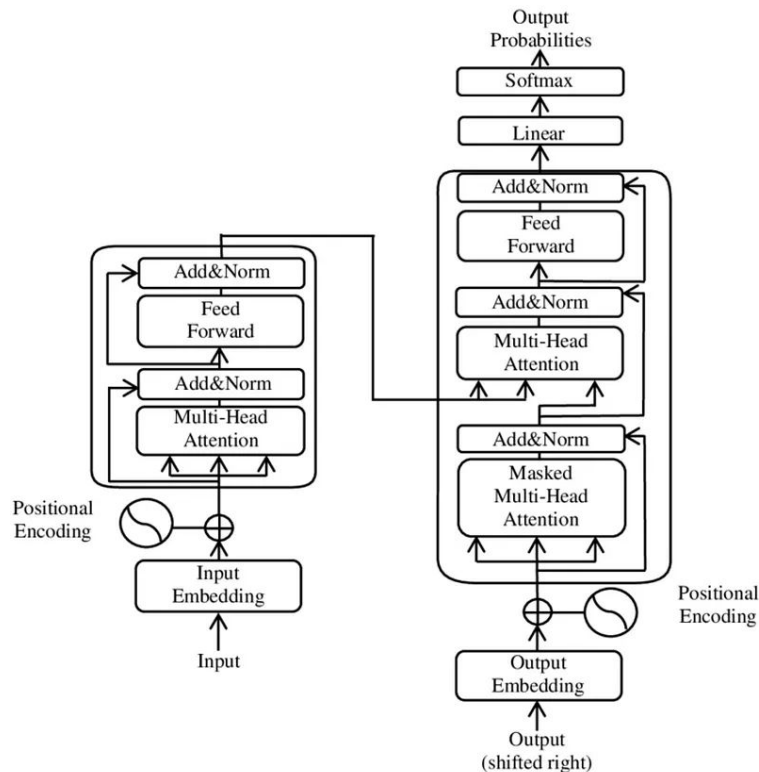



NA [A Vaswani, N Shazeer, N Parmar...](#) - Advances in neural ..., 2017 - [proceedings.neurips.cc](#)


... to attend to **all** positions in the decoder up to and including that position. **We need** to prevent ... **We** implement this inside of scaled dot-product **attention** by masking out (setting to $-\infty$) ...

☆ Save [Cite](#) Cited by 198960 [Related articles](#) [All 70 versions](#) [↔](#)

Transformers (Attention Mechanism)

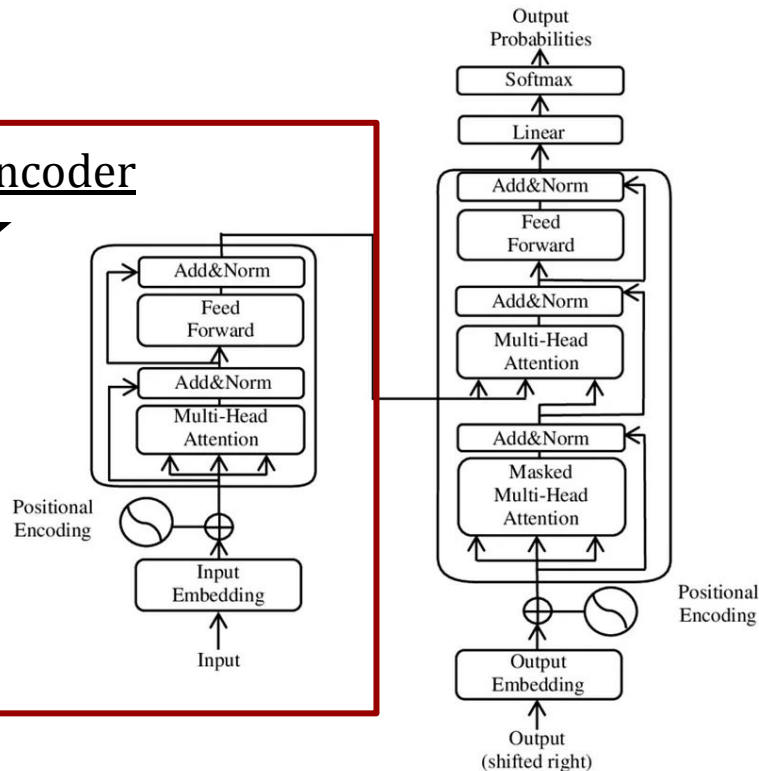


Transformers (Attention Mechanism)

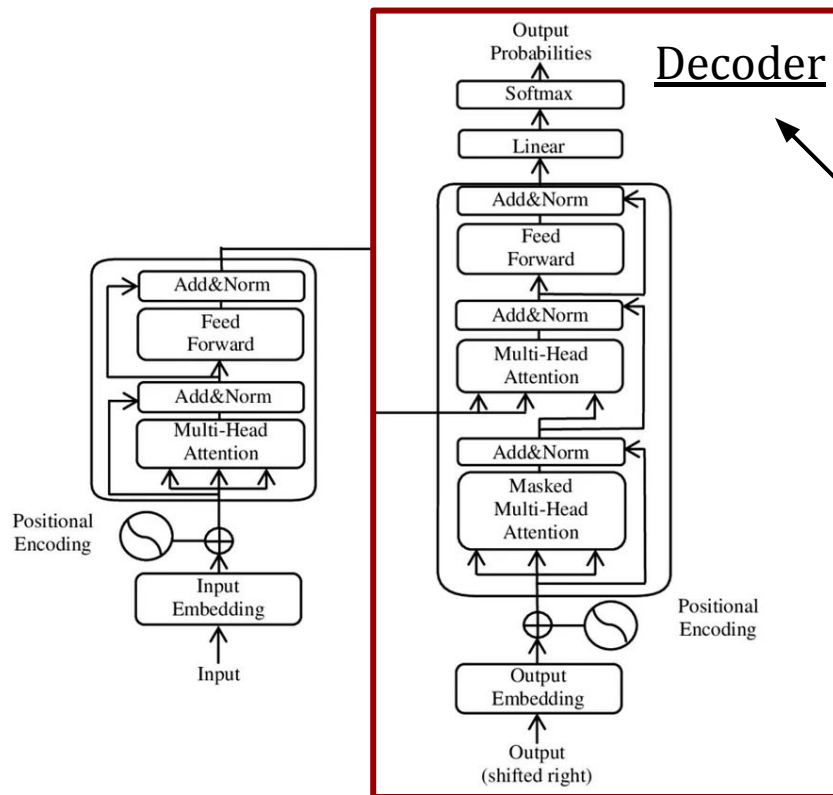



Encoder

An encoder transforms input data into a compressed representation while maintaining its meaning.



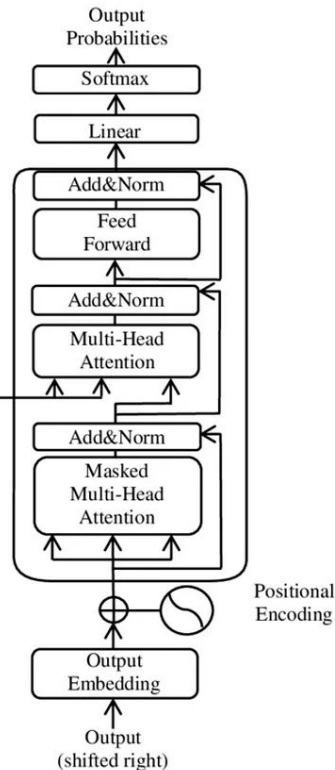
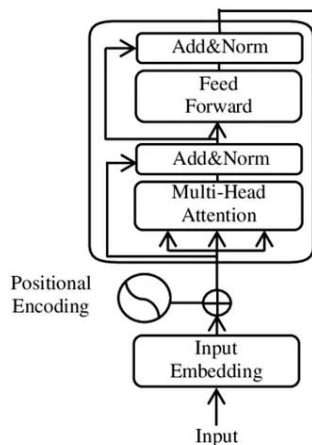
Transformers (Attention Mechanism)



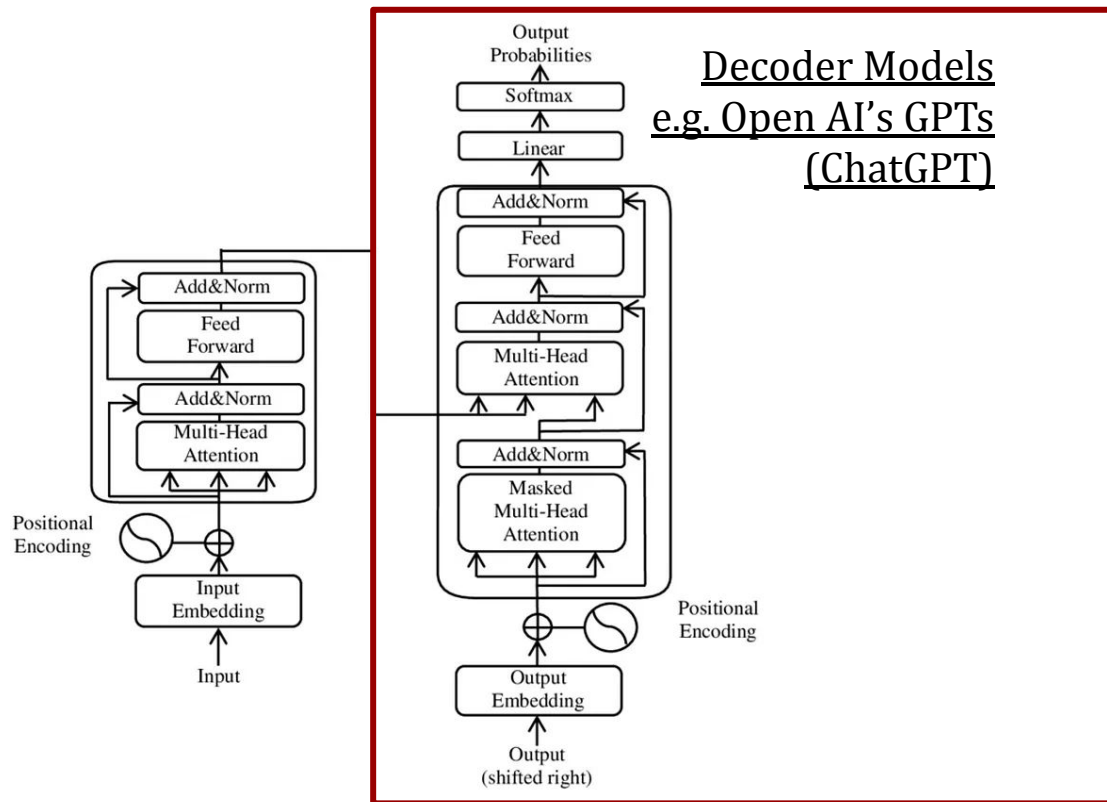
A decoder uses the representations of an encoder to generate new or reconstruct the desired output.

Transformers (Attention Mechanism)

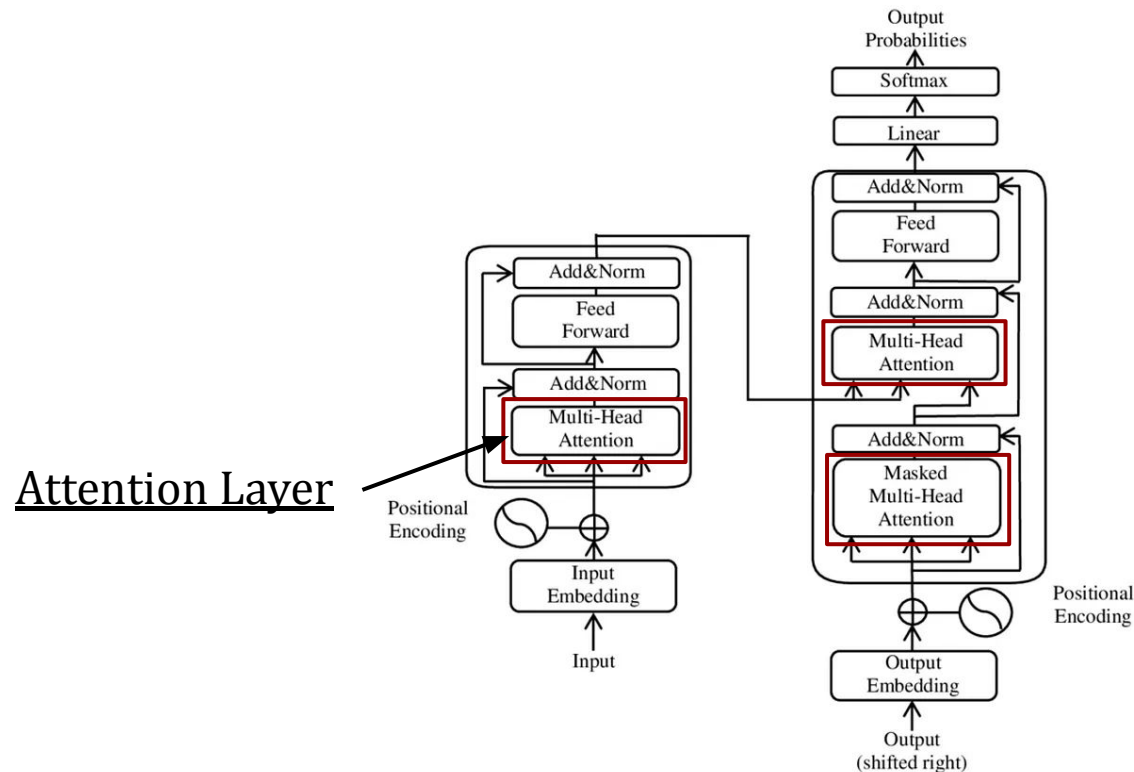
Encoder Models e.g., classifiers



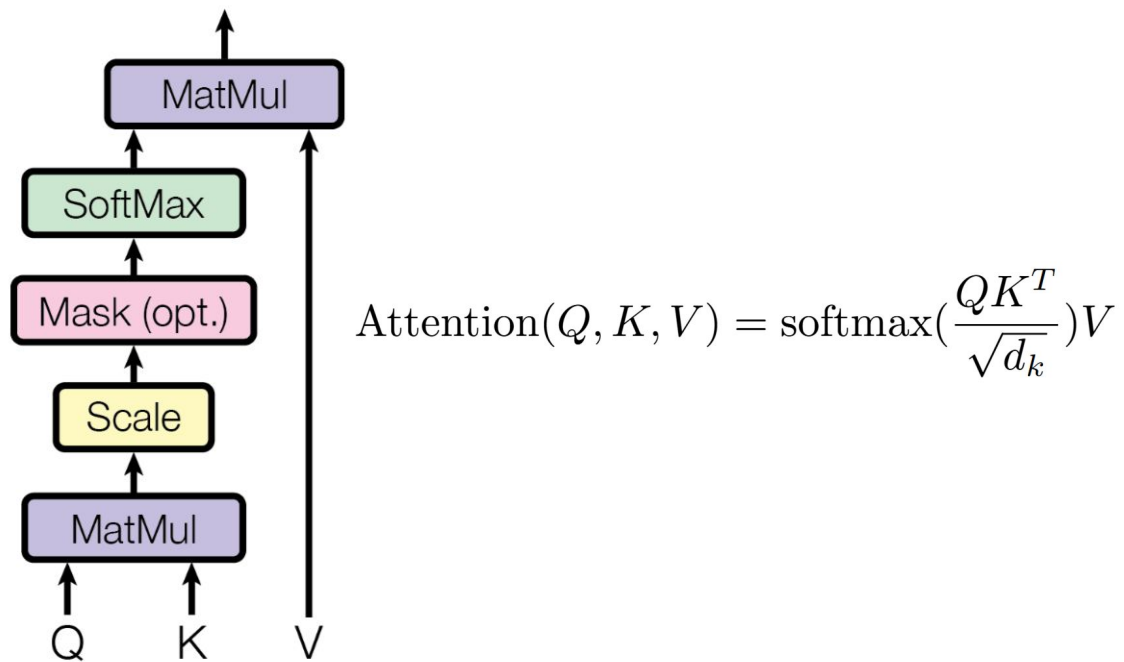
Transformers (Attention Mechanism)



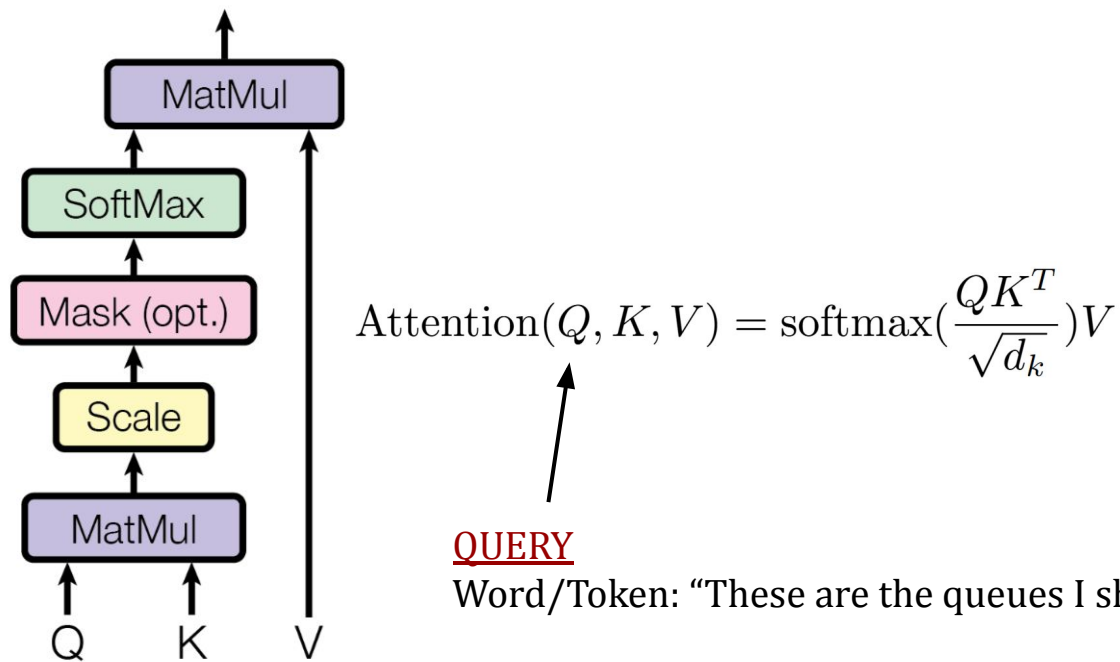
Transformers (Attention Mechanism)



Attention Layer




Attention Layer



QUERY

Word/Token: "These are the queues I should pay attention to."

Attention Layer



KEY

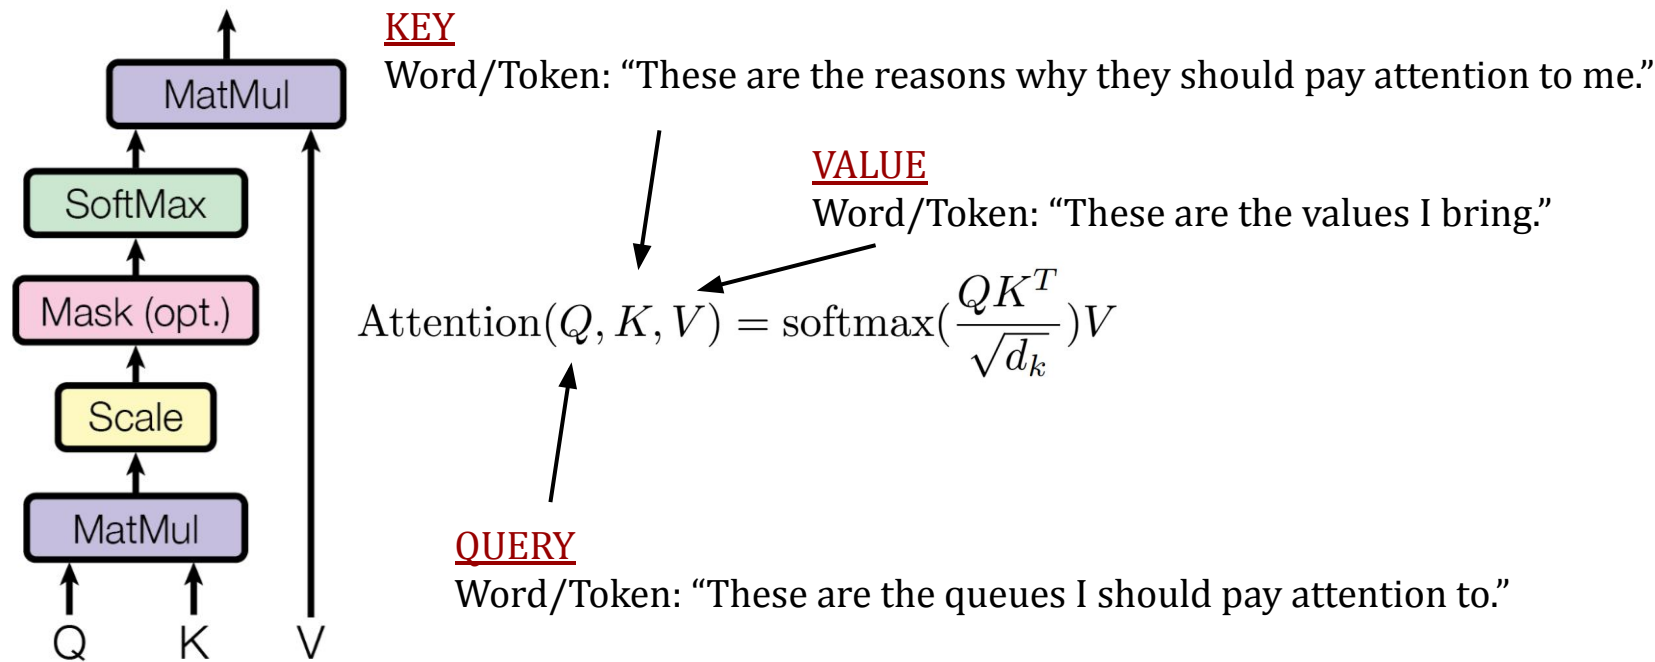
Word/Token: “These are the reasons why they should pay attention to me.”

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$


QUERY

Word/Token: “These are the queues I should pay attention to.”

Attention Layer




Attention Layer



Result: fixed representation matrices
(same as word2vec).

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

However, the non-linear interactions
of these matrices present
countless contexts trained in parallel,
not unique to words.



Pre-training Models: Self-supervision

Random
mask

A quick [MASK] fox jumps over the [MASK] dog.

Prediction

A quick brown fox jumps over the lazy dog.


Why are Transformers That Good?

- Attention Mechanism: Allows to model sequences (but RNNs can do that too!)
- Transformers can be massively parallelized. Ideal for GPUs!
- No memory loss as the RNNs. They take the entire text at the same time.
- And a few other tricks that we cannot cover...

Fine-tuning and Adaptation of Encoders

Encoder

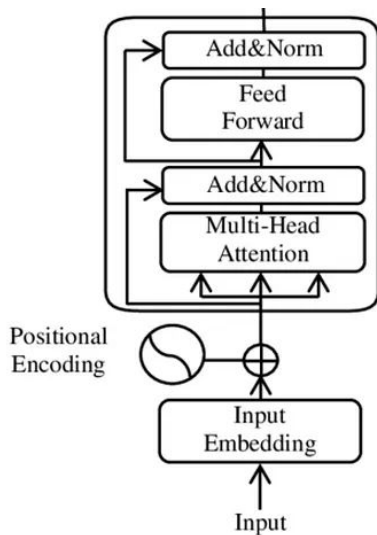
Pre-trained with massive amounts of text
using self-supervision



Fine-tuning and Adaptation of Encoders

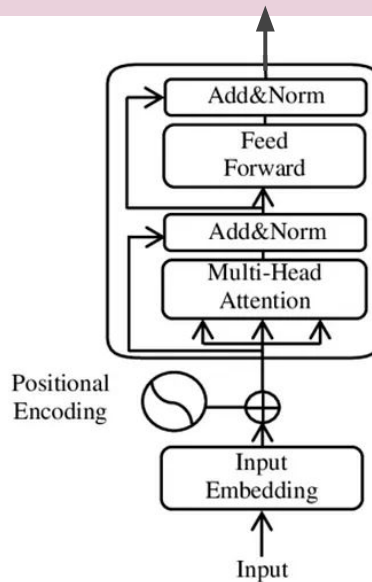
Encoder

Pre-trained with massive amounts of text
using self-supervision



Fine-tuning

Classification Layer
(predicting manually labeled text)



Lightweight adjustment of an
already language-aware models

Hugging Face: Transformers Repository



Hugging Face

Search models, dataset:

Models

Datasets

Spaces

Docs

Pricing



Log In

Sign Up

Main Tasks Libraries **Languages 1** Licenses
Other

Filter Languages by name

Reset Languages

- English ×
- Chinese
- French
- Spanish
- German
- Japanese
- Korean
- Italian
- Portuguese
- Russian
- Hindi
- Arabic
- Thai
- multilingual
- Turkish
- Indonesian
- Vietnamese
- Polish
- Dutch
- Romanian
- Swedish
- Ukrainian
- Czech
- Persian
- Finnish
- Bengali
- Nepali
- Malay
- Greek

Models 253,480

Filter by n:

Full-text search

Sort: Most downloads



sentence-transformers/all-MiniLM-L6-v2



Sentence Similarity • 22.7M • Updated Mar 6 • 118M • 4k



google/electra-base-discriminator

Updated Feb 29, 2024 • 58.5M • 66



google-bert/bert-base-uncased



Fill-Mask • 0.1B • Updated Feb 19, 2024 • 53.3M • 2.44k



FacebookAI/roberta-large



Fill-Mask • 0.4B • Updated Feb 19, 2024 • 20.4M • 250

Hugging Face: Applying and Fine-tuning a Model

```
import torch
from transformers import AutoModelForSequenceClassification
from transformers import AutoTokenizer

# Load tokenizer
tokenizer = AutoTokenizer.from_pretrained("luerhard/PopBERT")

# Load model
model = AutoModelForSequenceClassification.from_pretrained("luerhard/PopBERT")

# define text to be predicted
text = (
    "Das ist Klassenkampf von oben, das ist Klassenkampf im Interesse von "
    "Vermögenden und Besitzenden gegen die Mehrheit der Steuerzahlerinnen und "
    "Steuerzahler auf dieser Erde."
)

# encode text with tokenizer
encodings = tokenizer(text, return_tensors="pt")

# predict
with torch.inference_mode():
    out = model(**encodings)

# get probabilities
probs = torch.nn.functional.softmax(out.logits)
```



Hugging Face



luerhard / PopBERT

Hugging Face: Applying and Fine-tuning a Model

```
import torch
from transformers import AutoModelForSequenceClassification
from transformers import AutoTokenizer

# Load tokenizer
tokenizer = AutoTokenizer.from_pretrained("luerhard/PopBERT")

# Load model
model = AutoModelForSequenceClassification.from_pretrained("luerhard/PopBERT")

# define text to be predicted
text = (
    "Das ist Klassenkampf von oben, das ist Klassenkampf im Interesse von "
    "Vermögenden und Besitzenden gegen die Mehrheit der Steuerzahlerinnen und "
    "Steuerzahler auf dieser Erde."
)

# encode text with tokenizer
encodings = tokenizer(text, return_tensors="pt")

# predict
with torch.inference_mode():
    out = model(**encodings)

# get probabilities
probs = torch.nn.functional.softmax(out.logits)
```



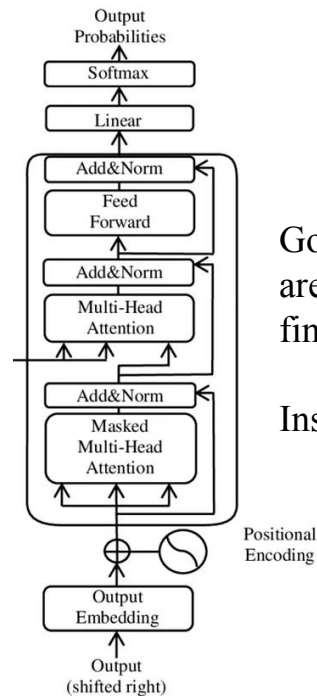
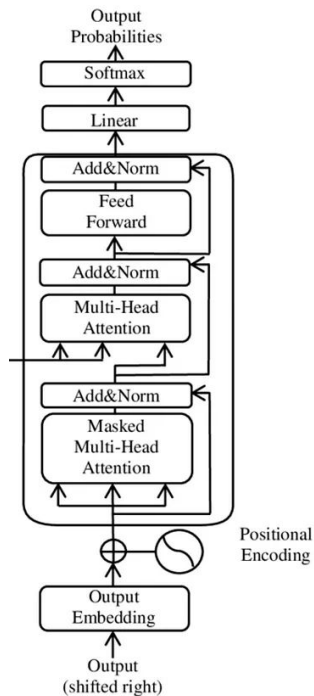
Hugging Face



luerhard/PopBERT

Check `fine_tuning.py`

Generative AI and Prompt Engineering



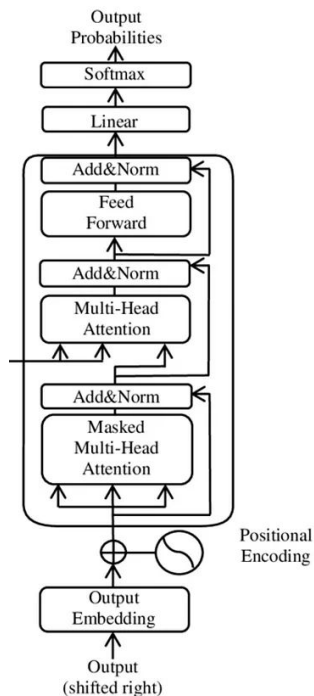
Good all-purpose generative models are too large and modular to be fine-tuned by a small group.

Instead...

Decoders/Generators

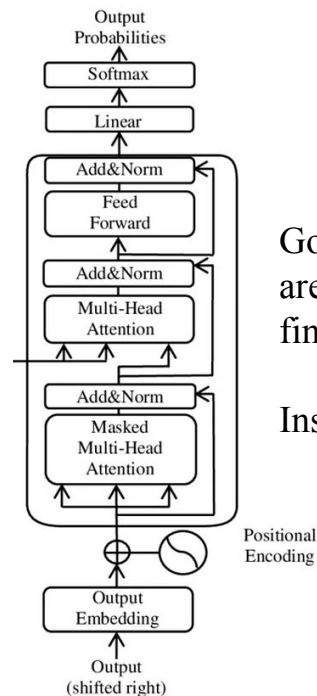
Text Generation / Question Answering

Generative AI and Prompt Engineering



Decoders/Generators

Text Generation / Question Answering



Good all-purpose generative models are too large and modular to be fine-tuned by a small group.

Instead...

Improving how we communicate with the model by providing more context.

Prompt Engineering

Prompt Engineering for Text Annotation

- Goal: Aim for **consistency**, **validity**, and **interpretability** of LLM-based annotations (e.g., sentiment, stance, ideology, populism, etc.)
- Steps:
 1. Specify task and schema clearly (**zero-shot vs. few-shot**)
 2. Control the **output format**
 3. Test **step by step reasoning** vs. **chain-of-thought suppression** (**temperature** = 0?)
 4. Test **robustness to prompts** (rephrase prompt, check quality against human annotators)
 5. Post-process, combine models/prompts (majority voting), **statistically validate**

Caution: small textual changes can affect construct validity and replicability.

Prompt Engineering for Text Annotation (API calls, Python)

```
import os
from litellm import completion

# --- Setup environment ---
os.environ["OPENAI_API_KEY"] = "sk-your-api-key-here" # <-- replace with your real key

# --- Input text ---
text = "The corrupt elite has betrayed the people!"

# --- Call ChatGPT via LiteLLM ---
response = completion(
    model="gpt-4o-mini",
    messages=[{"role": "user", "content": f"Label the following text as 'populism' or 'not populism':\n\n{text}"}])

# --- Output result ---
print(response["choices"][0]["message"]["content"])
```

Prompt Engineering for Text Annotation (API calls, [R](#))

```
library(ellmer)

chat <- chat_openai(model = "gpt-4.1")

text <- "The corrupt elite has betrayed the people!"

prompt <- paste0(
  "You are a classifier. Label the following text as **\"populism\"** or  

  **\"not populism\"** (just output the label, nothing else):\n\n", text
)

schema <- type_object(
  label = type_string()
)

res_struct <- chat$chat_structured(prompt, type = schema)

res_struct$label
```


Outline

- From word embeddings to Transformers
 - Encoding and Decoding: Understanding Modern Models
 - Fine-tuning and Adaptation
 - Generative AI and Prompt Engineering
-
- Tutorial: Sentiment and Topic Modeling in Parliamentary Speech
 - Discussion: The Promise and Pitfalls of NLP for Social Sciences

Tutorial

Link to notebook: https://github.com/boevkoski/intro_to_css_text_analysis_II

Discussion

Grimmer, Justin, and Brandon M. Stewart. "Text as data: The promise and pitfalls of automatic content analysis methods for political texts." Political analysis 21, no. 3 (2013): 267-297.

Discussion

Grimmer, Justin, and Brandon M. Stewart. "Text as data: The promise and pitfalls of automatic content analysis methods for political texts." Political analysis 21, no. 3 (2013): 267-297.

Table 1 Four principles of quantitative text analysis

-
- (1) All quantitative models of language are wrong—but some are useful.
 - (2) Quantitative methods for text amplify resources and augment humans.
 - (3) There is no globally best method for automated text analysis.
 - (4) Validate, Validate, Validate.
-

Next Class: Social network analysis from a sociological perspective II

- Statistical models in social network analysis and advanced topics
- Coding examples and exercises (with Mark)

Reading: *Snijders, Tom AB. "Statistical models for social networks." Annual review of sociology 37, no. 1 (2011): 131-153.*