

# Text Analysis II

## Introduction to Computational Social Science

DEPARTMENT OF NETWORK AND DATA SCIENCE (DNDS)

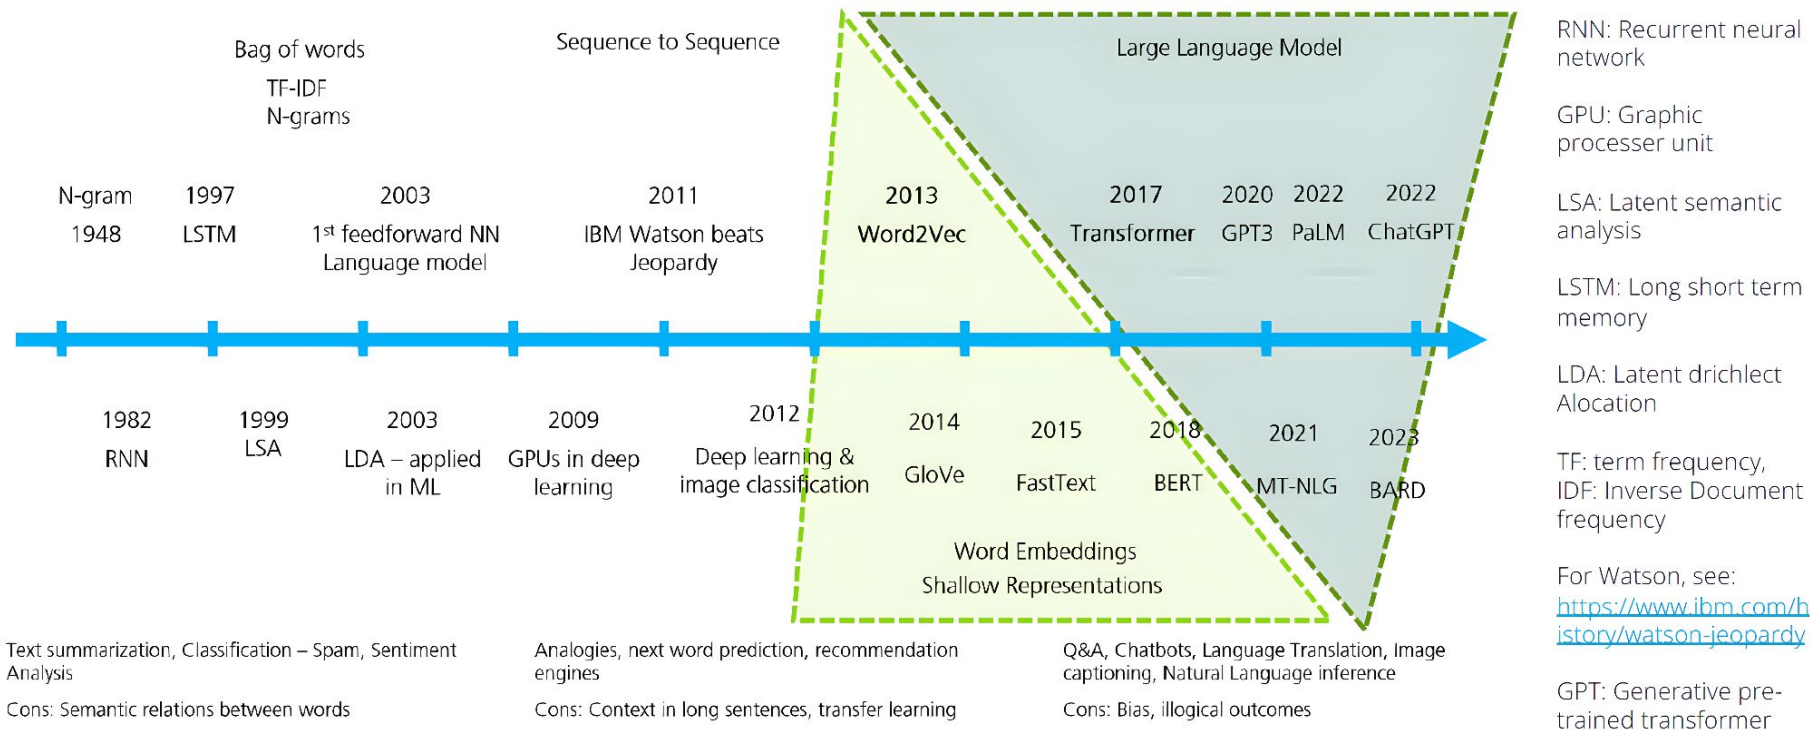
BOJAN EVKOSKI

OCT 2025

# Outline



- From word embeddings to Transformers
  - Encoding and Decoding: Understanding Modern Models
  - Fine-tuning and Adaptation
  - Generative AI and Prompt Engineering
- 
- Tutorial: Sentiment and Topic Modeling in Parliamentary Speech
  - Discussion: The Promise and Pitfalls of NLP for Social Sciences

# From word embeddings to Transformers




# Recurrent Neural Networks and LSTMs

one to one



# Recurrent Neural Networks and LSTMs

one to one



one to many






Image Captioning

many to one




Sentiment

many to many




Machine Translation

many to many





Named Entity  
Recognition

# Recurrent Neural Networks and LSTMs




**Issues:** slow to train; memory loss

# Word Embeddings





$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_t | w_{t+j})$$




$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

# Word Embeddings



$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_t | w_{t+j})$$



$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

**Issues:** aggregates contexts; single embedding for a unique word (e.g., Slack)



# Transformers (Attention Mechanism)

## Attention is all you need

[PDF] [neurips.cc](#)

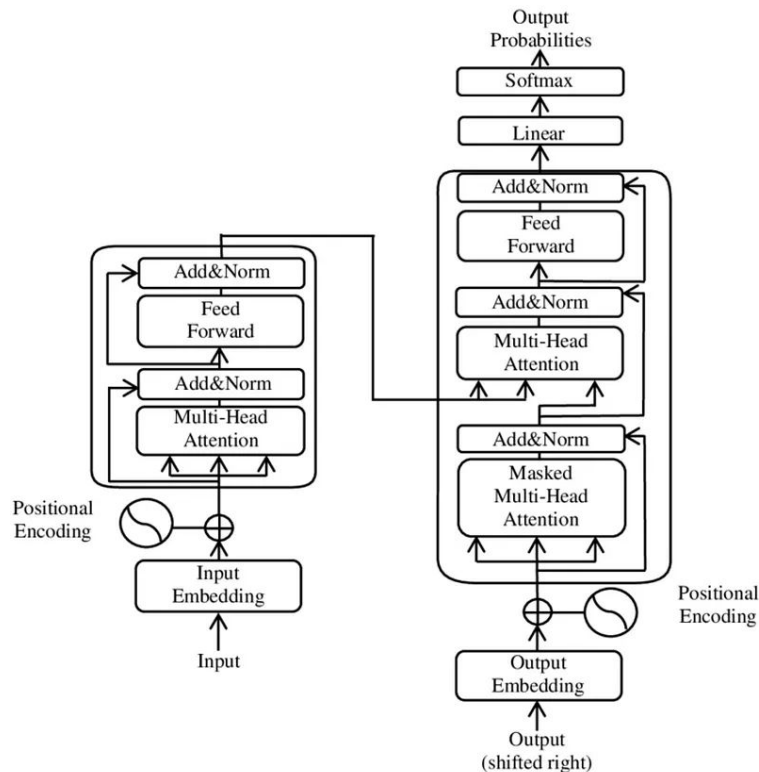



NA [A Vaswani](#), [N Shazeer](#), [N Parmar](#)... - Advances in neural ..., 2017 - [proceedings.neurips.cc](#)


... to attend to **all** positions in the decoder up to and including that position. **We need** to prevent ... **We** implement this inside of scaled dot-product **attention** by masking out (setting to  $-\infty$ ) ...

☆ Save [Cite](#) Cited by 198960 [Related articles](#) [All 70 versions](#) [↔](#)

# Transformers (Attention Mechanism)

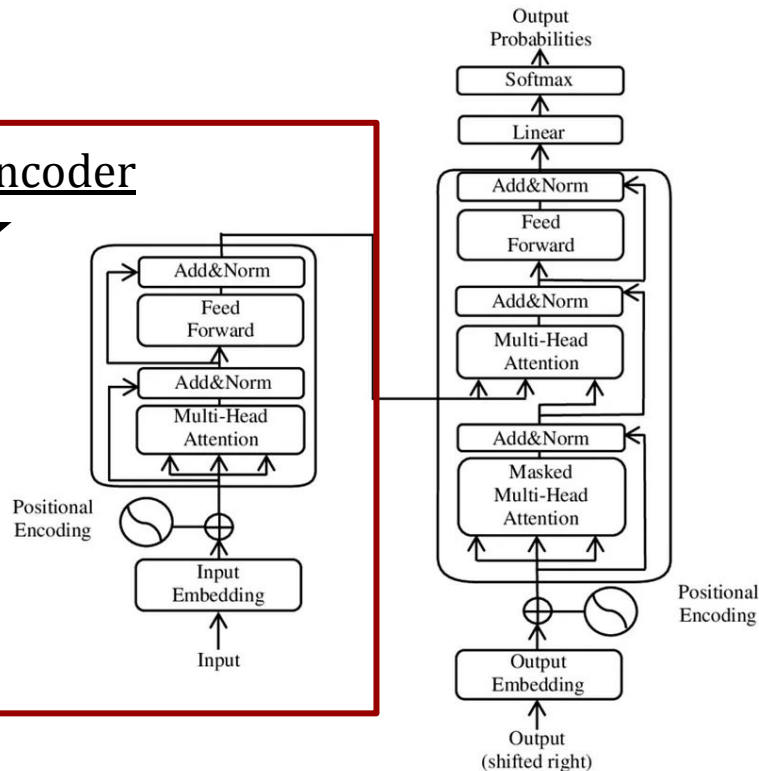


# Transformers (Attention Mechanism)

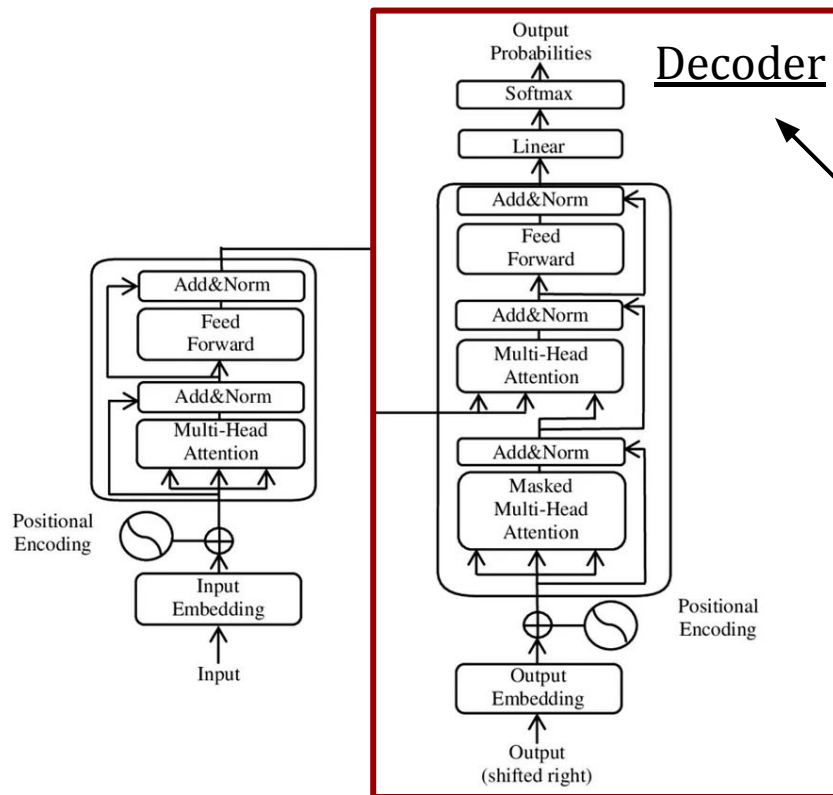



## Encoder

An encoder transforms input data into a compressed representation while maintaining its meaning.



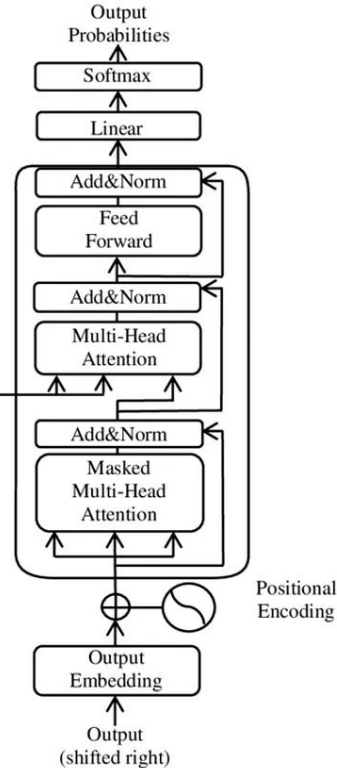

# Transformers (Attention Mechanism)



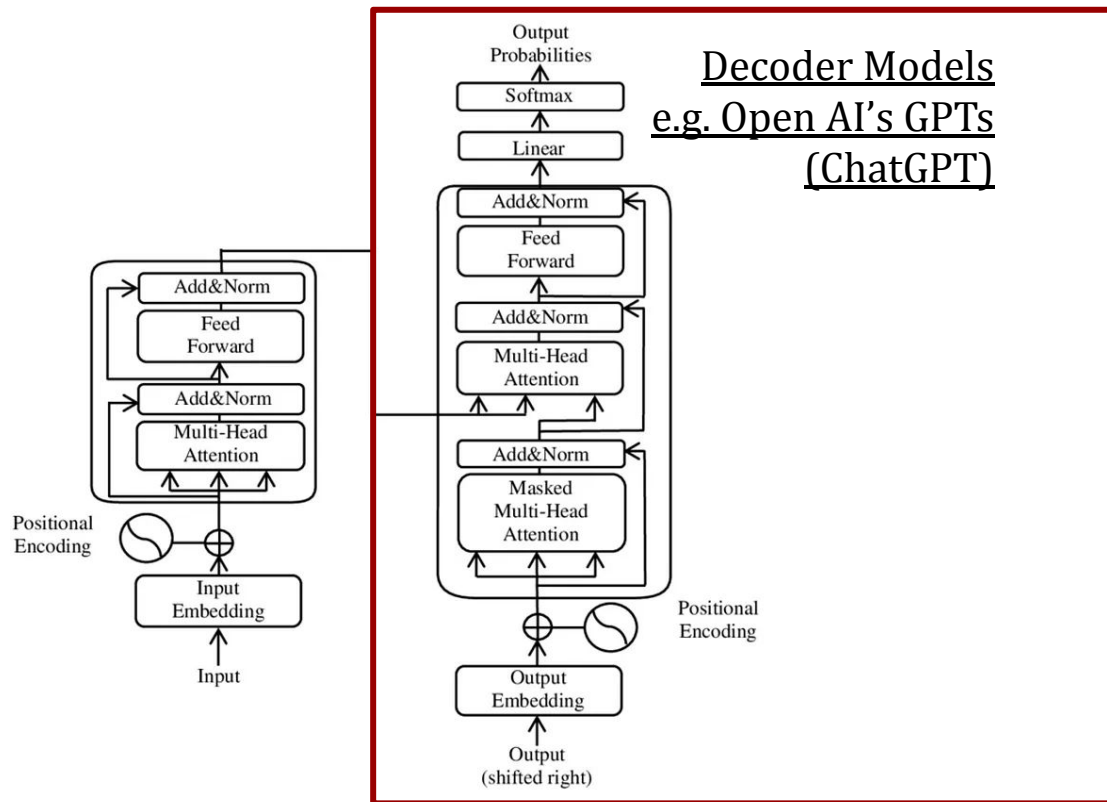
A decoder uses the representations of an encoder to generate new or reconstruct the desired output.

# Transformers (Attention Mechanism)

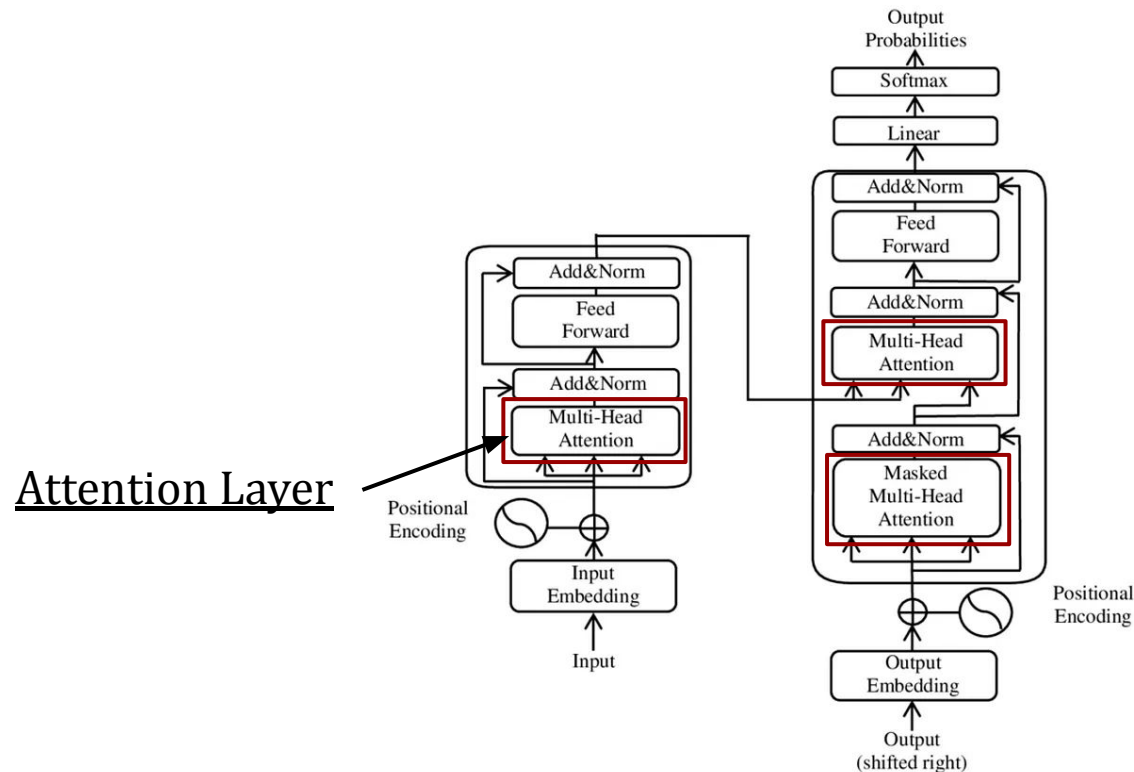
## Encoder Models e.g., classifiers



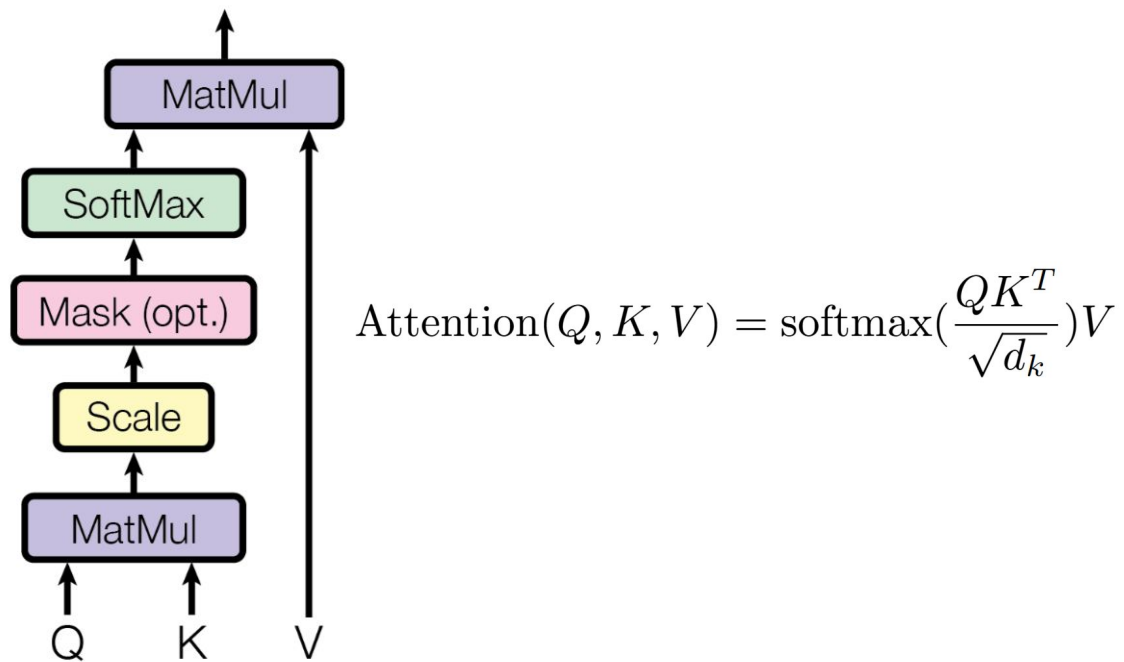
# Transformers (Attention Mechanism)



# Transformers (Attention Mechanism)

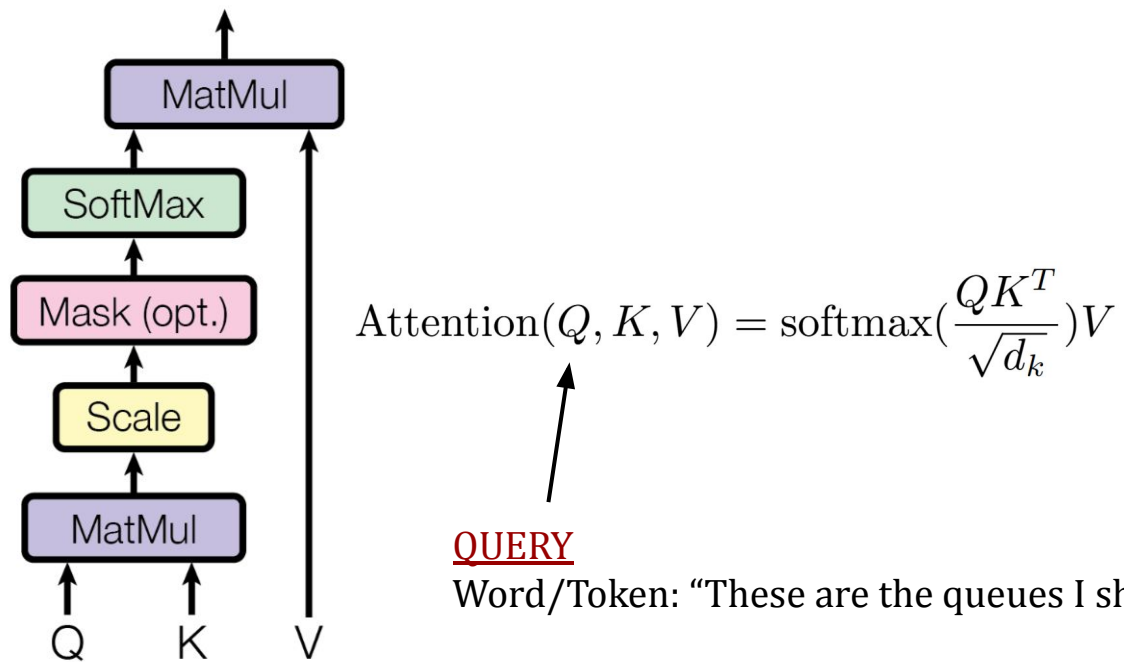


# Attention Layer






# Attention Layer



QUERY

Word/Token: "These are the queues I should pay attention to."

# Attention Layer



KEY

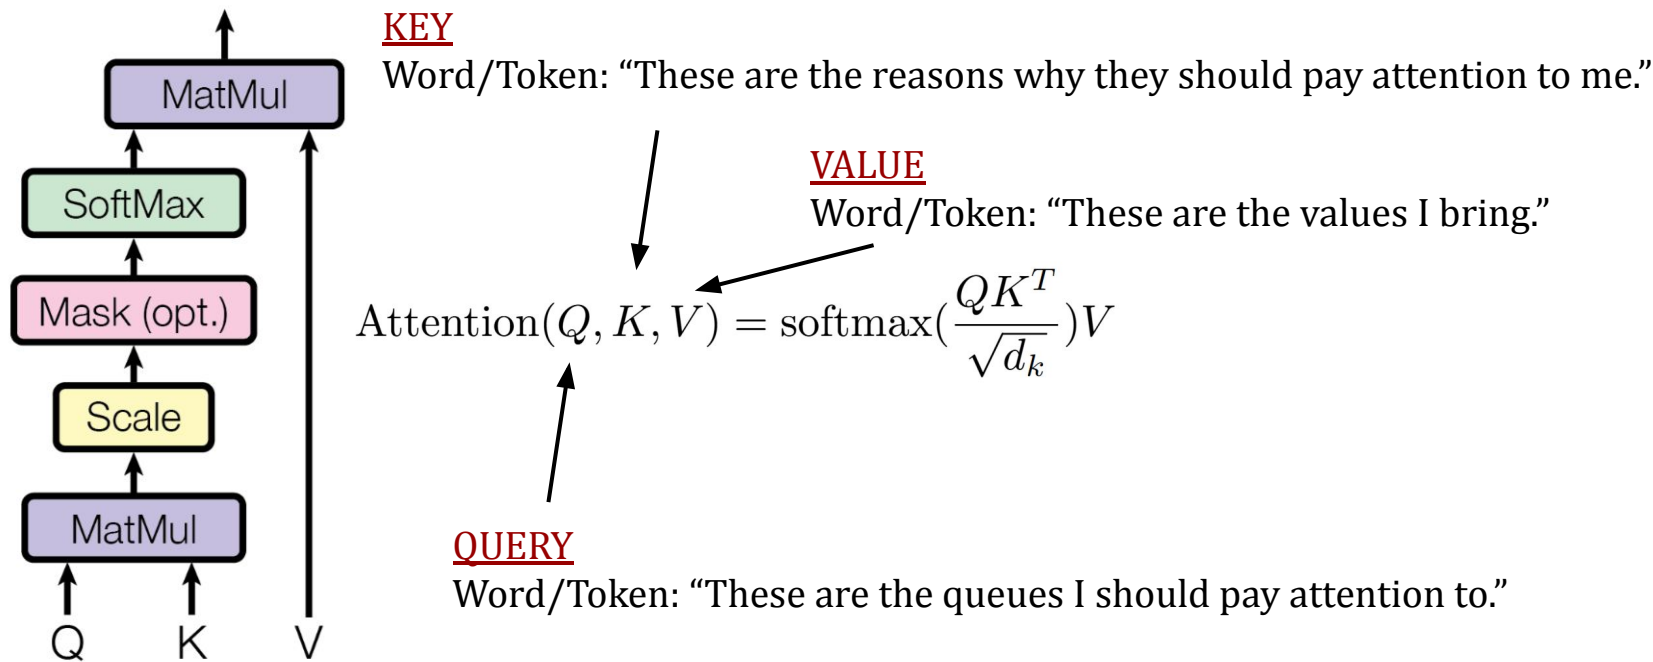
Word/Token: “These are the reasons why they should pay attention to me.”

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$


QUERY

Word/Token: “These are the queues I should pay attention to.”

# Attention Layer




# Attention Layer



**Result:** fixed representation matrices  
(same as word2vec).

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

However, the non-linear interactions  
of these matrices present  
**countless contexts trained in parallel**,  
not unique to words.



# Pre-training Models: Self-supervision

Random  
mask

A quick [MASK] fox jumps over the [MASK] dog.

Prediction

A quick brown fox jumps over the lazy dog.


# Why are Transformers That Good?

- Attention Mechanism: Allows to model sequences (but RNNs can do that too!)
- Transformers can be massively parallelized. Ideal for GPUs!
- No memory loss as the RNNs. They take the entire text at the same time.
- And a few other tricks that we cannot cover...

# Fine-tuning and Adaptation of Encoders

## Encoder


Pre-trained with massive amounts of text  
using self-supervision



# Fine-tuning and Adaptation of Encoders

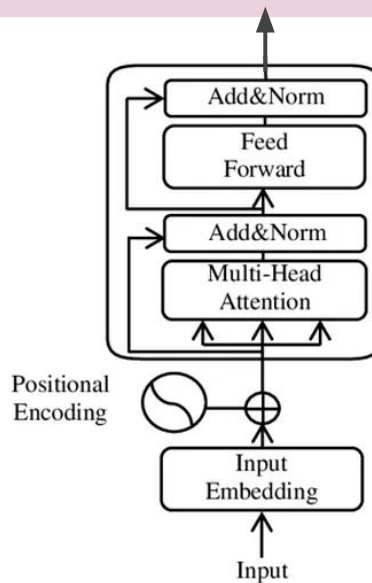
## Encoder

Pre-trained with massive amounts of text  
using self-supervision



## Fine-tuning

Classification Layer  
(predicting manually labeled text)



Lightweight Adjustment of  
already language-aware models



# Hugging Face: Transformers Repository



Hugging Face

Search models, dataset:

Models

Datasets

Spaces

Docs

Pricing



Log In

Sign Up

Main Tasks Libraries **Languages 1** Licenses  
Other

Filter Languages by name

Reset Languages

English × Chinese French Spanish  
German Japanese Korean Italian  
Portuguese Russian Hindi Arabic  
Thai multilingual Turkish  
Indonesian Vietnamese Polish  
Dutch Romanian Swedish  
Ukrainian Czech Persian Finnish  
Bengali Nepali Malay Greek

Models 253,480

Filter by n:

Full-text search

Sort: Most downloads



sentence-transformers/all-MiniLM-L6-v2



Sentence Similarity • 22.7M • Updated Mar 6 • 118M • 4k



google/electra-base-discriminator

Updated Feb 29, 2024 • 58.5M • 66



google-bert/bert-base-uncased



Fill-Mask • 0.1B • Updated Feb 19, 2024 • 53.3M • 2.44k



FacebookAI/roberta-large



Fill-Mask • 0.4B • Updated Feb 19, 2024 • 20.4M • 250

# Hugging Face: Applying and Fine-tuning a Model

```
import torch
from transformers import AutoModelForSequenceClassification
from transformers import AutoTokenizer

# Load tokenizer
tokenizer = AutoTokenizer.from_pretrained("luerhard/PopBERT")

# Load model
model = AutoModelForSequenceClassification.from_pretrained("luerhard/PopBERT")

# define text to be predicted
text = (
    "Das ist Klassenkampf von oben, das ist Klassenkampf im Interesse von "
    "Vermögenden und Besitzenden gegen die Mehrheit der Steuerzahlerinnen und "
    "Steuerzahler auf dieser Erde."
)

# encode text with tokenizer
encodings = tokenizer(text, return_tensors="pt")

# predict
with torch.inference_mode():
    out = model(**encodings)

# get probabilities
probs = torch.nn.functional.softmax(out.logits)
```



**Hugging Face**



**luerhard / PopBERT**

# Hugging Face: Applying and Fine-tuning a Model

```
import torch
from transformers import AutoModelForSequenceClassification
from transformers import AutoTokenizer

# Load tokenizer
tokenizer = AutoTokenizer.from_pretrained("luerhard/PopBERT")

# Load model
model = AutoModelForSequenceClassification.from_pretrained("luerhard/PopBERT")

# define text to be predicted
text = (
    "Das ist Klassenkampf von oben, das ist Klassenkampf im Interesse von "
    "Vermögenden und Besitzenden gegen die Mehrheit der Steuerzahlerinnen und "
    "Steuerzahler auf dieser Erde."
)

# encode text with tokenizer
encodings = tokenizer(text, return_tensors="pt")

# predict
with torch.inference_mode():
    out = model(**encodings)

# get probabilities
probs = torch.nn.functional.sigmoid(out.logits)
```



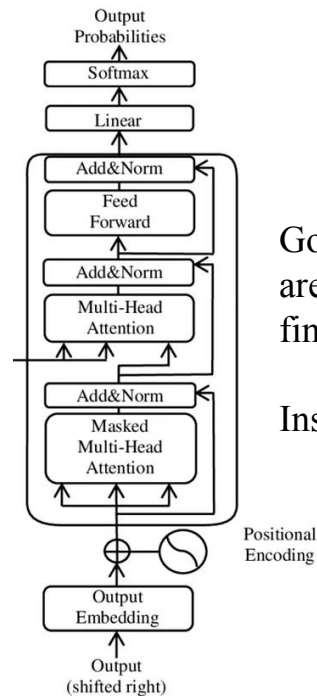
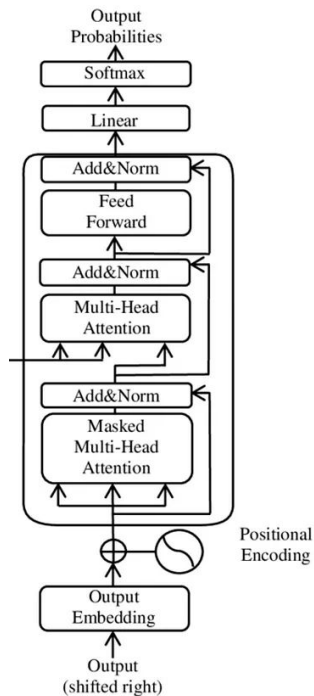
**Hugging Face**



**luerhard/PopBERT**

Check `fine_tuning.py`

# Generative AI and Prompt Engineering




Good all-purpose generative models are too large and modular to be fine-tuned by a small group.

Instead...

## Decoders/Generators


Text Generation / Question Answering

# Generative AI and Prompt Engineering



## Decoders/Generators

Text Generation / Question Answering



Good all-purpose generative models are too large and modular to be fine-tuned by a small group.

Instead...

Improving how we communicate with the model by providing more context.

## Prompt Engineering

# Prompt Engineering for Text Annotation

- Goal: Improve consistency, validity, and interpretability of LLM-based annotations (e.g., sentiment, stance, ideology, populism, etc.)
- Steps:
  1. Specify task and schema clearly (**zero-shot vs. few-shot**)
  2. Control the **output format**
  3. Test **step by step reasoning** vs. **chain-of-thought suppression** (**temperature** = 0?)
  4. Test Prompt Robustness (rephrase prompt, check quality against human annotators)
  5. Post-process, combine models/prompts (majority voting), statistically validate

Caution: small textual changes can affect construct validity and replicability.

# Prompt Engineering for Text Annotation (API calls, Python)

```
import os
from litellm import completion

# --- Setup environment ---
os.environ["OPENAI_API_KEY"] = "sk-your-api-key-here" # <-- replace with your real key

# --- Input text ---
text = "The corrupt elite has betrayed the people!"

# --- Call ChatGPT via LiteLLM ---
response = completion(
    model="gpt-4o-mini",
    messages=[{"role": "user", "content": f"Label the following text as 'populism' or 'not populism':\n\n{text}"}])

# --- Output result ---
print(response["choices"][0]["message"]["content"])
```

# Prompt Engineering for Text Annotation (API calls, [R](#))

```
library(ellmer)

chat <- chat_openai(model = "gpt-4.1")

text <- "The corrupt elite has betrayed the people!"

prompt <- paste0(
  "You are a classifier. Label the following text as **\"populism\"** or  

  **\"not populism\"** (just output the label, nothing else):\n\n", text
)

schema <- type_object(
  label = type_string()
)

res_struct <- chat$chat_structured(prompt, type = schema)

res_struct$label
```



# Outline

- From word embeddings to Transformers
  - Encoding and Decoding: Understanding Modern Models
  - Fine-tuning and Adaptation
  - Generative AI and Prompt Engineering
- 
- Tutorial: Sentiment and Topic Modeling in Parliamentary Speech
  - Discussion: The Promise and Pitfalls of NLP for Social Sciences

# Tutorial

Link to notebook: [https://github.com/boevkoski/intro\\_to\\_css\\_text\\_analysis\\_II](https://github.com/boevkoski/intro_to_css_text_analysis_II)

# Discussion

*Grimmer, Justin, and Brandon M. Stewart. "Text as data: The promise and pitfalls of automatic content analysis methods for political texts." Political analysis 21, no. 3 (2013): 267-297.*

# Discussion

*Grimmer, Justin, and Brandon M. Stewart. "Text as data: The promise and pitfalls of automatic content analysis methods for political texts." Political analysis 21, no. 3 (2013): 267-297.*

**Table 1** Four principles of quantitative text analysis

- 
- (1) All quantitative models of language are wrong—but some are useful.
  - (2) Quantitative methods for text amplify resources and augment humans.
  - (3) There is no globally best method for automated text analysis.
  - (4) Validate, Validate, Validate.
-